

Database

Database is a collection of related data serves as a container for

- `tables`
- `indexes`
- `views`
- `etc.`

and other database `objects`.

Create new database

Syntax:

```
CREATE DATABASE db_name
WITH
    [OWNER = role]
    [TEMPLATE = template]
    [ENCODING = encoding]
    [LC_COLLATE = collate]
    [LC_TYPE = ctype]
    [TABLESPACE = tablespace_name]
    [ALLOW_CONNECTIONS = true|false]
    [CONNECTION LIMIT = max_concurrent_connection]
    [IS_TEMPLATE = true|false]
```

- **OWNER:** Assign a role(user) that will be the owner of the database. the owner have the highest level of control over the database.

The owner has `full privileges` over the `database`, including the ability to `drop` it, `alter` it, and `grant/revoke` privileges to other `users`.

```
OWNER = myuser

-- DEFAULT USER IS postgres
```

- **TEMPLATE:** Specify the `template` database for the new database. postgres use the `[template1]` database as default template database.
 - PostgreSQL comes with two default templates:
 - `template0`: A minimal database with no user-created objects.
 - `template1`: A more commonly used template that may include objects and customizations.

```
TEMPLATE=template1
```

NOTE:

The database created from a template inherits all objects (**tables**, **functions**, **etc.**) in the template. However, certain templates, like **template0**, are clean and have no **user-defined** objects.

- **ENCODING:**

- Defines the character encoding for the database.
- This setting determines how **text** is stored in the **database**.
- Common encodings include **UTF8** (recommended for most applications).
- Common values: **UTF-8**, **LATIN1**, **SQL_ASCII**, etc.

```
ENCODING= 'UTF-8'
```

- **LC_COLLATE:**

- Specifies the **collation** order to use for string **sorting** and **comparison**.
- it determines the rules for character ordering in the database.
- **Collations** are based on **locale** settings and can differ for different **languages** and **regions**.
- Common values: **en_US.UTF-8**, **fr_FR.UTF-8**, etc

```
LC_COLLATE = 'en_US.UTF-8'
```

NOTE:

You cannot change these properties after the database is created, so they should be set carefully.

- **LC_TYPE:** Defines the **character** classification and **case conversion** behavior, such as upper and lower case conversions. like **LC_COLLATE**, this is local-based.

Common values: **en_US.UTF-8**, **fr_FR.UTF-8**, etc

```
LC_TYPE = 'en_US.UTF-8'
```

NOTE:

You cannot change these properties after the database is created, so they should be set carefully.

- **TABLESPACE:** Specifies the tablespace where the **database** should reside. A **tablespace** is a **storage** location on the **disk** where the database **files** are stored.
 - If not specified, **PostgreSQL** uses the default tablespace.
 - The default tablespace is **PostgreSQL** is **pg_default**, which maps to **/data** directory in **PostgreSQL**.
 - **PostgreSQL** also has a second default tablespace called **pg_global**, which stores global data.

```
TABLESPACE=my_tablespace

-- if you want to print available tablespace use below commands.
SELECT spcname from pg_tablespace;

-- show the physical location of tablespace
show data_directory;

-- To list all tables that are stored in the pg_default tablespace
SELECT tabelname from pg_tables;
```

- **ALLOW_CONNECTIONS:** Determines whether the database should allow connections. Setting this to false will prevent anyone from connecting to the database, but the database will still exist for administrative purposes.

Possible values: **TRUE** | **FALSE**

```
ALLOW_CONNECTIONS = true
```

- **CONNECTION LIMIT:** Defines the maximum number of concurrent connections allowed to the database. A **-1** value (or omitting this option) means unlimited connections.

```
CONNECTION LIMIT = 100
```

- **IS_TEMPLATE:** Specifies whether the database should be treated as a template for creating new databases.

Possible values: **TRUE** | **FALSE**

```
IS_TEMPLATE=true
```

Example:

```
CREATE DATABASE blogs
  OWNER = pkuser
  TEMPLATE = template1
  ENCODING = 'UTF8'
  LC_COLLATE = 'en_US.UTF-8'
  LC_CTYPE = 'en_US.UTF-8'
  TABLESPACE = my_tablespace
  ALLOW_CONNECTIONS = true
  CONNECTION LIMIT = 100
  IS_TEMPLATE = false;

-- OR
-- rest using default values of each parameter
CREATE DATABASE blogs;

-- retrieve the database names from the `pg_database`
SELECT datname from pg_database;

-- list all database in [psql]
\l

-- connect to created or any database
\c db_name
```

Alter database

ALTER DATABASE statement allow you to carry the following action on the database.

- Change the attributes of the database.
- Rename the database.
- Change the owner of the database.
- Change the default tablespace of a database.
- Change the session default for a non-runtime configuration variable for a database.

Changing attributes of a database

```
-- syntax
ALTER DATABASE name WITH option;

-- option can be;
-- IS_TEMPLATE
-- CONNECTION LIMIT
-- ALLOW_CONNECTIONS
```

```
-- Only superusers or database owner can change these settings;
```

Rename database:

```
-- syntax
ALTER DATABASE db_name
RENAME TO new_db_name;

-- It is not possible to rename the current database.
-- Only superusers and database owners with [CREATEDB] privilege can rename the
database.
```

Change the owner of the database:

```
-- syntax
ALTER DATABASE db_name
OWNER TO new_owner | current_user | session_user;

-- to check current user or session user run below command
SELECT current_user; -- Returns the current role executing the query
SELECT session_user; -- Returns the role that authenticated the session
SELECT user; -- Equivalent to SELECT CURRENT_USER;

-- session information: pg_stat_activity
SELECT username, application_name, client_addr, backend_start, state
FROM ps_stat_activity
WHERE pid=pg_backend_pid();

-- In PostgreSQL, [ pg_backend_pid()] is a function that returns the [ process ID
(PID)] of the current backend process.

-- how to terminate the current session
SELECT pg_terminate_backend(pg_backend_pid());
```

Change the default tablespace of a database:

```
-- syntax
ALTER DATABASE db_name
SET TABLESPACE new_tablespace;

-- To set the new tablespace, the tablespace needs to be empty and there is a
connection to the database.
-- Superusers and database owners can change the default tablespace of the
database
```

Change session defaults for run-time configuration variables:

Whenever you connect to a **database**, PostgreSQL loads the **configuration variables** from the **postgresql.conf** file and uses these variables by default

```
-- syntax
ALTER DATABASE database_name
SET configuration_parameter = value;

-- check the current setting from [pg_settings]
SELECT name, setting FROM pg_settings;

-- change configuration variables for current session
SET <configuration_parameter> TO <value>;

-- examples
SET work_mem TO '64MB'; -- (Memory used for sorting operations)
SET search_path TO my_schema, public; -- (Schema search order)
SET timezone TO 'UTC'; -- (Time zone setting)
SET statement_timeout TO '5min'; -- (Maximum execution time for a statement)
SET log_statement TO 'ddl'; -- (Logging level for SQL statements)
SET default_transaction_isolation TO 'READ COMMITTED'; -- (Transaction isolation level)

-- Reset configuration variable to default values
RESET <<variable_name>>
-- OR
RESET ALL; -- reset all

-- exmple reset work_mem only
RESET work_mem;
```

Examples:

```
-- create database
CREATE DATABASE testdb2;

-- rename to testhrdb
ALTER DATABASE testdb2
RENAME TO testhrdb;

-- change owner postgres to hr
ALTER DATABASE testhrdb
OWNER TO hr;

-- change the default tablespace of the testhrdb from pg_default to hr_default
ALTER DATABASE testhrdb
SET TABLESPACE hr_default;
```

```
-- set escape_string_warning configuration variable to off by using the following
statement:
ALTER DATABASE testhrdb
SET escape_string_warning = off;
```

Drop database

The **DROP DATABASE** statement deletes a database from a PostgreSQL server.

```
-- syntax
DROP DATABASE [IF EXISTS] database_name
[WITH (FORCE)]

-- The FORCE option will attempt to terminate all existing connections to the
target database.
```

NOTE:

- The **DROP DATABASE** statement deletes the database from both **catalog entry** and **data directory**.
- Since PostgreSQL does not allow you to **roll back** this operation, you should use it with caution.
- To execute the **DROP DATABASE** statement, you need to be the database **owner**.

Examples:

```
-- Create some database
CREATE DATABASE hr;
CREATE DATABASE test;

-- Drop the hr database
DROP DATABASE hr;

-- Removing a non-existing database example (IF EXISTS WILL CHECK THE DATABASE
THEN DELETE IF EXISTS OTHERWISE DO NOTHING)
DROP DATABASE IF EXISTS non_existing_database;

-- Drop a database that has active connections example
DROP DATABASE test WITH (FORCE)
```

Rename database

```
-- Create database bots
CREATE DATABASE bots;

-- RENAME TO robots
```

```
ALTER DATABASE bots
RENAME TO robots;
```

Copy database within the same server

You want to copy a PostgreSQL database within a database server for testing purpose.

```
-- syntax
CREATE DATABASE targetDb
WITH TEMPLATE sourceDb;

-- example
CREATE DATABASE dvdrental_test
WITH TEMPLATE dvdrental;
```

Copy database from one server to another:

```
-- Step-1: dump the source database into a file.
pg_dump -U postgres -d sourcedb -f sourcedb.sql

-- Step-2: create new database
CREATE DATABASE demo;

-- Step-3: restore the dump file
psql -U postgres -d demo -f sourcedb.sql
```

How to Get Sizes of Database Objects in PostgreSQL

- Use the `pg_size_pretty()` function to format the size.
- Use the `pg_relation_size()` function to get the size of a table.
- Use the `pg_total_relation_size()` function to get the total size of a table.
- Use the `pg_database_size()` function to get the size of a database.
- Use the `pg_indexes_size()` function to get the size of an index.
- Use the `pg_total_index_size()` function to get the size of all indexes on a table.
- Use the `pg_tablespace_size()` function to get the size of a tablespace.
- Use the `pg_column_size()` function to obtain the size of a column of a specific type.

Examples:

```
-- Getting table sizes:
select pg_relation_size('actor');

-- The pg_size_pretty() function formats a number using bytes, kB, MB, GB, or TB
appropriately. For example:
SELECT
    pg_size_pretty (pg_relation_size('actor')) size;
```



```
-- To get the total size of a table
SELECT
    pg_size_pretty (
        pg_total_relation_size ('actor')
    ) size;

-- the following query returns the top 5 biggest tables in the dvdrental database
SELECT
    relname AS "relation",
    pg_size_pretty (
        pg_total_relation_size (C .oid)
    ) AS "total_size"
FROM
    pg_class C
LEFT JOIN pg_namespace N ON (N.oid = C .relnamespace)
WHERE
    nspname NOT IN (
        'pg_catalog',
        'information_schema'
    )
AND C .relkind <> 'i'
AND nspname !~ '^pg_toast'
ORDER BY
    pg_total_relation_size (C .oid) DESC
LIMIT 5;

-- Getting database size
SELECT
    pg_size_pretty (
        pg_database_size ('dvdrental')
    ) size;

-- Getting index sizes
SELECT
    pg_size_pretty (pg_indexes_size('actor')) size;

-- Getting tablespace sizes
SELECT
    pg_size_pretty (
        pg_tablespace_size ('pg_default')
    ) size;

-- Getting PostgreSQL value sizes
SELECT
    pg_column_size(5 :: smallint) smallint_size,
    pg_column_size(5 :: int) int_size,
    pg_column_size(5 :: bigint) bigint_size;
```